

HTML_2023_HW4

tags: Personal

Discuss with anonymous and TAs.

Ref: [TA Poy HTML 2023 HW4](<https://hackmd.io/@Poy/SJVAvg0Qn>)

More about Regularization

P1

In order to find the minimum w , we can set the derivative of the optimal solution to be zero, hence, we have

$$\begin{aligned}
0 &= \frac{\partial}{\partial w} \left(\frac{1}{N} \sum_{n=1}^N (wx_n - y_n)^2 + \frac{\lambda}{N} w^2 \right) \\
&= \frac{2}{N} \sum_{n=1}^N x_n (wx_n - y_n) + \frac{2w\lambda}{N} \\
&= \frac{2w}{N} \sum_{n=1}^N x_n^2 - \frac{2}{N} \sum_{n=1}^N x_n y_n + \frac{2w\lambda}{N} \\
&= \frac{2w}{N} \left(\sum_{n=1}^N x_n^2 + \lambda \right) - \frac{2}{N} \sum_{n=1}^N x_n y_n \\
\frac{2w}{N} \left(\sum_{n=1}^N x_n^2 + \lambda \right) &= \frac{2}{N} \sum_{n=1}^N x_n y_n \\
w \left(\sum_{n=1}^N x_n^2 + \lambda \right) &= \sum_{n=1}^N x_n y_n \\
w^* &= \frac{\sum_{n=1}^N x_n y_n}{\sum_{n=1}^N x_n^2 + \lambda} \\
\text{or we can use } w^* &= (x^T x + \lambda I)^{-1} x^T y \\
&= \left(\sum_{n=1}^N x_n^2 + \lambda \right)^{-1} \cdot \sum_{n=1}^N x_n y_n \\
&= \frac{\sum_{n=1}^N x_n y_n}{\sum_{n=1}^N x_n^2 + \lambda} \\
\therefore C &= (w^*)^2 \\
\therefore C &= \left(\frac{\sum_{n=1}^N x_n y_n}{\sum_{n=1}^N x_n^2 + \lambda} \right)^2
\end{aligned}$$

As a result, we should choose $[a]$ as our solution.

P2

Since the L2-regularized linear regression in the \mathcal{Z} -space is equivalent to regularized linear regression in the \mathcal{X} -space, we have

$$\begin{aligned}
&\min_{\tilde{w} \in \mathbb{R}^{d+1}} \frac{1}{N} \sum_{n=1}^N (\tilde{w}^T \Phi(x_n) - y_n)^2 + \frac{\lambda}{N} (\tilde{w}^T \tilde{w}) \\
&= \min_{\tilde{w} \in \mathbb{R}^{d+1}} \frac{1}{N} \sum_{n=1}^N (w^T x_n - y_n)^2 + \frac{\lambda}{N} \Omega(w) \\
&\therefore \tilde{w}^T \Phi(x_n) = \tilde{w} \Gamma^{-1} (x - u) = \tilde{w} \Gamma^{-1} x = w^T x_n, \tilde{w}^T \tilde{w} = \Omega(w)
\end{aligned}$$

Since Γ is a diagonal matrix, we know that the inverse of diagonal matrix equals to itself, which means $\Gamma^T = \Gamma$ then we have

$$\begin{aligned}
w^T &= \tilde{w} \Gamma^{-1} \\
\tilde{w}^T &= w^T \Gamma \\
\tilde{w} &= (w^T \Gamma)^T = \Gamma^T w \\
\therefore \Omega(w) &= \tilde{w}^T \tilde{w} = w^T \Gamma \Gamma w = w^T \Gamma^2 w
\end{aligned}$$

As a result, we should choose $[b]$ as our solution.

P3

$$\min_w \frac{1}{N} \sum_{n=1}^N \text{err}_{\text{smooth}}(w, x_n, y_n) = \min_w \frac{1}{N} \sum_{n=1}^N \text{err}(w, x_n, y_n) + \frac{\lambda}{N} \sum_{n=1}^N \Omega(w, x_n)$$

$$D_{KL}(P\|Q) = \sum_x P(x) \ln \frac{P(x)}{Q(x)} = \sum_x \frac{1}{2} \ln \frac{\frac{1}{2}}{\text{err}(w, x, y)}$$

$$\begin{aligned}
w_{lr} &= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n D_{KL}(P(Y_i) \| P(\hat{Y}_i)) \\
&= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n y_i \ln \frac{y_i}{p_i} + (1 - y_i) \ln \frac{(1 - y_i)}{(1 - p_i)} \\
&= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n y_i (\ln y_i - \ln p_i) + (1 - y_i) (\ln(1 - y_i) - \ln(1 - p_i)) \\
&= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n -y_i \ln p_i - (1 - y_i) \ln(1 - p_i) + (y_i \ln y_i + (1 - y_i) \ln(1 - y_i)) \\
&= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n y_i \ln p_i + (1 - y_i) \ln(1 - p_i) \\
&= \underset{w}{\operatorname{argmin}} \sum_{i=1}^n H(P(Y_i) \| P(\hat{Y}_i)) \\
\therefore \Omega(w, x) &= D_{KL}(P_u \| P_h)
\end{aligned}$$

As a result, we should choose $[a]$ as our solution.

Validation

P4

Since constant hypothesis $h(x) = w_0$, we only need to take y into the consideration. By selecting two examples for predicting another one example, we can obtain the equation

$$\begin{aligned}
E_{loocv} &= \frac{1}{3} \left(\left(y_1 - \frac{y_2 + 1}{2} \right)^2 + \left(y_2 - \frac{y_1 + 1}{2} \right)^2 + \left(1 - \frac{y_1 + y_2}{2} \right)^2 \right) \\
&= \frac{1}{2} (y_1^2 + y_2^2 - y_1 y_2 - y_1 - y_2 + 1)
\end{aligned}$$

Since we are finding $E_{loocv} \leq \frac{1}{3}$, we can rewrite the equation to be

$$\begin{aligned} & \left(\frac{x+1}{2} - y\right)^2 + \left(\frac{y+1}{2} - x\right)^2 + \left(\frac{x+y}{2} - 1\right)^2 = 1 \\ \Rightarrow & \frac{1}{2}(x^2 + y^2 - xy - x - y + 1) = \frac{1}{3} \end{aligned}$$

Drawing the equation with GeoGebra, we can simply get the area, which is $\approx 2.4183 \approx \frac{\pi}{3\sqrt{3}}$



As a result, we should choose $[b]$ as our solution.

P5

$$\begin{aligned}
& \because E_{out}(h) = \mathcal{E}_{x \sim P} [h(x) \neq f(x)] \\
& \therefore \text{Variance}_{\mathcal{D}_{val} \sim \mathcal{P}^K} [E_{val}(h)] \\
&= \text{Variance}_{(x,y) \sim \mathcal{P}} \left[\frac{1}{K} \sum_{n=1}^K \text{err}(h(x_n), y_n) \right] \\
&= \frac{1}{K^2} \text{Variance}_{(x,y) \sim \mathcal{P}} \left[\sum_{n=1}^K \text{err}(h(x_n), y_n) \right] \\
&= \frac{1}{K^2} \text{Variance}_{(x,y) \sim \mathcal{P}} [\text{err}(h(x_1), y_1) + \text{err}(h(x_2), y_2) + \dots + \text{err}(h(x_K), y_K)] \\
&= \frac{1}{K^2} (K \cdot \text{Variance}_{(x,y) \sim \mathcal{P}} [\text{err}(h(x), y)]) \\
&= \frac{1}{K} \text{Variance}_{(x,y) \sim \mathcal{P}} [\text{err}(h(x), y)] \\
&\therefore \square = \frac{1}{K}
\end{aligned}$$

Since the examples (x, y) are generated from the i.i.d. distribution, the covariance between examples should be zero.

$$\begin{aligned}
& \because \text{Covariance} [\text{err}(h(x_i), y_i), \text{err}(h(x_j), y_j)] = 0, \forall i, j \\
& \therefore \text{Variance} (A + B) = \text{Variance} (A) + \text{Variance} (B) + 2 \cdot \text{Covariance} (A, B) \\
& \quad = \text{Variance} (A) + \text{Variance} (B) + 0
\end{aligned}$$

As a result, we should choose $[d]$ as our solution.

P6

Since single data set in a binary classification has N positive examples and N negative examples, we can perform leave-one-out validation by selecting one positive or one negative example into the validation data set.

Once we select a positive example as a validation data set, there are $N - 1$ positive examples and N negative examples in the training data set, and the binary classification algorithm $A_{majority}$ will return negative to the training data set for the reason that it predicts the majority class. However, $A_{majority}$ will return positive to the validation data set since there is only one positive example, and vice versa.

As a result, we know that $A_{majority}$ will always return a different classification between the training data set and the validation data set, which means $E_{loocv}(A_{majority}) = 1$.

As a result, we should choose $[d]$ as our solution.

P7

The threshold of the decision stump model can be represented as

$$\theta \in \{-1\} \cup \left\{ \frac{x_i + x_{i+1}}{2} : 1 \leq i \leq N-1 \text{ and } x_i \neq x_{i+1} \right\}$$

Since x is generated by a uniform distribution in $[-1, +1]$, and $y = \text{sign}(x)$, it's trivial that $E_{in} = 0$ when $\theta = 0$, and the threshold θ with the lowest E_{in} in the decision stump model will be trained in the middle of the smallest positive value and the largest negative value. Therefore, the trained $\theta = \frac{x_j + x_{j+1}}{2}$, where x_j is the largest negative value, and x_{j+1} is the smallest positive value.

If the x_j is selected as the validation data set, then $\theta = \frac{x_{j-1} + x_{j+1}}{2}$, where x_{j-1} is the second largest negative value. Both θ and the x_j are between x_{j-1} and x_{j+1} , then we have

- if $\theta \geq x_j \Rightarrow h(x_j) = -1$, which means the validation data set will be classified correctly.
- if $\theta < x_j \Rightarrow h(x_j) = +1$, which means the validation data set will be classified incorrectly.

And vice versa, if x_{j+1} is selected as the validation data set, then $\theta = \frac{x_j + x_{j+2}}{2}$, where x_{j+2} is the second smallest positive value. Both θ and the x_{j+1} are between x_j and x_{j+2} , then we have

- if $\theta \geq x_{j+1} \Rightarrow h(x_{j+1}) = -1$, which means the validation data set will be classified incorrectly.
- if $\theta < x_{j+1} \Rightarrow h(x_{j+1}) = +1$, which means the validation data set will be classified correctly.

Since only x_j and x_{j+1} will be classified incorrectly by the threshold θ , the tightest upper bound of the leave-one-out validation error to the decision stump model is equal to $\max(E_{loocv}) = 2/N$.

As a result, we should choose $[c]$ as our solution.

Support Vector Machine

P8

Since the examples are ordered, the hard-margin SVM without transformation can be considered as a decision stump model when the separation line is orthogonal to the coordinate of the examples. It's trivial that the perfect separation line must be the middle of x_M and x_{M+1} , therefore, the hypothesis can be represented as

$$g_{svm}(x) = \text{sign}(w^T x + b) = 0 = x - \frac{x_M + x_{M+1}}{2}$$

Therefore, the largest margin is equal to $\frac{1}{2}(x_{M+1} - x_M)$

As a result, we should choose $[e]$ as our solution.

P9

Based on the subjections of $\min(w, b) = \frac{1}{2} w^T w$, we have

$$\begin{aligned} (w^T x_n + b) &\geq 1 && \text{for } y_n = +1 \\ -(w^T x_n + b) &\geq 1126 && \text{for } y_n = -1 \end{aligned}$$

$$y = \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}^T \begin{bmatrix} 0 & 4 \\ 2 & 0 \\ -1 & 0 \\ 0 & 0 \end{bmatrix} + b = \begin{bmatrix} +1 \\ -1 \\ +1 \\ +1 \end{bmatrix} \Rightarrow \begin{bmatrix} (4w_2 + b) \geq 1 \\ -(2w_1 + b) \leq 1126 \\ (-w_1 + b) \geq 1 \\ (+b) \geq 1 \end{bmatrix}$$

By drawing all data on the graph and from the subjections above, we know that the ratio of positive margin versus negative margin can represent as $r_m = \frac{y_{n=+1}}{y_{n=-1}} = \frac{1}{1126}$, hence, the hyperplane can be represented by the thinnest gap $-(w^T [2, 0] + b = 1126)$ and $(w^T [0, 0] + b = 1)$, which implies that $w^T x + b = 0 \Rightarrow -\frac{1127}{2}x + 1 = 0$, which means $w = (-\frac{1127}{2}, 0), b = 1$.

As a result, we should choose $[e]$ as our solution.

P10

Since $\alpha = 1, b = 0, h(x) = \text{sign}(\sum_{n=1}^N y_n K(x_n, x))$, then we have $E_{in}(h) = \frac{1}{N} \sum_{i=1}^N [h(x_i) \neq y_i]$. If $E_{in}(\hat{h}) = 0$, all predictions should be classified correctly, and there is no classification error, which means

$$\left[\text{sign} \left(\sum_{n=1}^N y_n K(x_n, x_i) \right) = y_i \right], \forall i \xrightarrow{\text{rewrite}} \left[\left(\sum_{n=1}^N y_n K(x_n, x_i) \right) \cdot y_i > 0 \right], \forall i$$

To sum up the equation above, we have

$$\begin{aligned} \sum_{i=1}^N \left(y_i \sum_{n=1}^N y_n K(x_n, x_i) \right) &= [y_1^2 + y_1 (y_2 K(x_2, x_1) + \cdots + y_N K(x_N, x_1))] \\ &\quad + [y_2^2 + y_2 (y_1 K(x_1, x_2) + \cdots + y_N K(x_N, x_2))] \\ &\quad + \cdots \\ &\quad + [y_N^2 + y_N (y_1 K(x_1, x_N) + \cdots + y_{N-1} K(x_{N-1}, x_N))] > 0 \end{aligned}$$

Since $\gamma > 0, \|x_n \neq x_m\| \geq \epsilon, \forall n \neq m$, we have

$$\begin{aligned} K(x_n, x_m) &= \exp(-\gamma \|x_n - x_m\|^2) \leq \exp(-\gamma \epsilon^2) := \zeta \\ \text{if } n = i, K(x_n, x_i) &= \exp(-\gamma \|x_n - x_i\|^2) = \exp(-\gamma \cdot 0) = 0 \end{aligned}$$

Taking ζ back to the equation as the upper bound, then we get

$$\begin{aligned}
0 &< \sum_{i=1}^N \left(y_i \sum_{n=1}^N y_n K(x_n, x_i) \right) < \sum_{i=1}^N \left(y_i \sum_{n=1}^N y_n \exp(-\gamma \epsilon^2) \right) \\
&= [y_1^2 + y_2^2 + \dots + y_N^2] \\
&+ \underbrace{\left[y_1 \left(\underbrace{y_2 \exp(-\gamma \epsilon^2) + \dots + y_N \exp(-\gamma \epsilon^2)}_{N-1} \right) + \dots + y_N \left(\underbrace{y_1 \exp(-\gamma \epsilon^2) + \dots + y_{N-1} \exp(-\gamma \epsilon^2)}_{N-1} \right) \right]}_N
\end{aligned}$$

And we know that

$$\begin{aligned}
\forall n = m, y_m = (+1) \vee (-1) &\Rightarrow y_i^2 = 1, \forall i \\
\forall n \neq m, -1 < y_n y_m < 1 &\Rightarrow \text{assumed that all } y_n y_m = -1
\end{aligned}$$

Then, we have

$$\begin{aligned}
0 &< \sum_{i=1}^N \left(y_i \sum_{n=1}^N y_n \exp(-\gamma \epsilon^2) \right) \\
&< \sum_{i=1}^N y_i^2 + \sum_{i=1}^N (-1) \times (N-1) \exp(-\gamma \epsilon^2) \\
&= N + N \times [-(N-1) \exp(-\gamma \epsilon^2)] \\
&= N - N(N-1) \exp(-\gamma \epsilon^2) \\
1 &> (N-1) \exp(-\gamma \epsilon^2) \\
\frac{1}{N-1} &> \exp(-\gamma \epsilon^2) \\
\ln(N-1) &< \gamma \epsilon^2 \\
\gamma &> \frac{\ln(N-1)}{\epsilon^2}
\end{aligned}$$

As a result, we should choose $[d]$ as our solution.

P11

Applying the feature transform ϕ to the Gaussian kernel, then we have

$$\begin{aligned}
\|\phi(x) - \phi(x')\|^2 &= \langle \phi(x) - \phi(x'), \phi(x) - \phi(x') \rangle \\
&= \langle \phi(x), \phi(x) \rangle - 2 \langle \phi(x), \phi(x') \rangle + \langle \phi(x'), \phi(x') \rangle \\
&= \phi(x)^T \phi(x) - \phi(x)^T \phi(x') - \phi(x')^T \phi(x) + \phi(x')^T \phi(x') \\
&= K(x, x) - 2K(x, x') + K(x', x') \\
&= 2 - 2 \exp(-\gamma \|x - x'\|^2) \\
&\because \exp(-\gamma \|x - x'\|^2) \in (0, 1] \\
&\therefore 2 > \|\phi(x) - \phi(x')\|^2 \\
\sqrt{2} \approx 1.5 &> \|\phi(x) - \phi(x')\|
\end{aligned}$$

Therefore, the tightest upper bound for the distance in the \mathcal{Z} -space is 1.5.

As a result, we should choose $[d]$ as our solution.

Experiments with Regularized Logistic Regression

P12

Since L2-regularized logistic regression is in liblinear solves

$$w = \min_w \frac{1}{2} w^T w + C \sum_{i=1}^N \log(1 + \exp(-y_i w^T x_i))$$

and we solve

$$w_\lambda = \operatorname{argmin}_w \frac{\lambda}{N} \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \ln(1 + \exp(-y_n w^T \Phi_4(x_n)))$$

Then we know that

$$\frac{1}{2} = \frac{\lambda}{N}, C = \frac{1}{N} \Rightarrow C = \frac{1}{2\lambda}$$

The parameter `-c cost` : set the parameter C (default 1) should be set to $C = \frac{1}{2\lambda}$.

```
1 import numpy as np
2 from itertools import combinations_with_replacement
3 from liblinear.liblinearutil import *
4
5 def transformation(args, x):
6     x_origin = x[:, 1:] # remove x_{n=0}=1
7     transformed_x = x.copy()
8     combination_tuples = list()
9     iterable = list(i for i in range(1, args['dimension']))
10    for r in range(1, args['Q']):
11        combination_tuples.extend(list(combinations_with_replacement(iterable, r+1)))
12
13    for combination in range(len(combination_tuples)):
14        x_temp = 1
15        for j in combination_tuples[combination]:
16            x_temp = x_origin[:,j-1].reshape(-1, 1) * x_temp
17        transformed_x = np.hstack((
18            transformed_x, x_temp
19        ))
20    return transformed_x
```

```

1 def read_file(args, filename):
2     data = np.loadtxt(filename, dtype=float)
3     x = data[:, :-1]
4     x = np.c_[np.ones(len(x)), x] # x_{n=0}=1
5     y = data[:, -1]
6     args['dimension'] = x.shape[1]
7     return x, y

```

```

1 def main(args):
2     x_train, y_train = read_file(args, args['filename_train'])
3     x_test, y_test = read_file(args, args['filename_test'])
4     x_train = transformation(args, x_train)
5     x_test = transformation(args, x_test)
6
7     E_out = list()
8     for lamb in args['lambda']:
9         C = 1/(2*lamb)
10        prob = problem(y_train, x_train)
11        param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))
12        m = train(prob, param)
13        p_label, p_acc, p_val = predict(y_test, x_test, m)
14        E_out.append(round(np.mean(y_test != p_label), 6))
15
16    min_E_out = min(E_out)
17    min_E_out_index = [i for i, v in enumerate(E_out) if v == min_E_out]
18    print("E_out: ", E_out)
19    print("min E_out: ", min(E_out))
20    print("Choice: ", chr(97+max(min_E_out_index)))

```

```

1 if __name__ == '__main__':
2     args = {
3         'dimension': 0,
4         'filename_test': "hw4_test.dat",
5         'filename_train': "hw4_train.dat",
6         'lambda': [10**(-6), 10**(-3), 10**(0), 10**3, 10**6],
7         'Q': 4,
8     }
9
10    main(args)

```

```

1 NumbaDeprecationWarning: The 'nopython' keyword argument was not supplied to the 'numl
2 def csr_to_problem_jit(l, x_val, x_ind, x_rowptr, prob_val, prob_ind, prob_rowptr):
3 Accuracy = 77.4% (387/500) (classification)
4 Accuracy = 82.2% (411/500) (classification)
5 Accuracy = 84.6% (423/500) (classification)
6 Accuracy = 85.8% (429/500) (classification)
7 Accuracy = 81.2% (406/500) (classification)
8 E_out: [0.226, 0.178, 0.154, 0.142, 0.188]
9 min E_out: 0.142
10 Choice: d

```

Therefore, the best $\log_{10}(\lambda^*) = 3$

As a result, we should choose $[d]$ as our solution.

P13

We use the same code above, then update the `main` function. The updated functions show below.

```

1 def main(args):
2     x_train, y_train = read_file(args, args['filename_train'])
3     x_test, y_test = read_file(args, args['filename_test'])
4     x_train = transformation(args, x_train)
5     x_test = transformation(args, x_test)
6
7     E_in = list()
8     for lamb in args['lambda']:
9         C = 1/(2*lamb)
10        prob = problem(y_train, x_train)
11        param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))
12        m = train(prob, param)
13        p_label, p_acc, p_val = predict(y_train, x_train, m)
14        E_in.append(round(np.mean(y_train != p_label), 6))
15
16    min_E_in = min(E_in)
17    min_E_in_index = [i for i, v in enumerate(E_in) if v == min_E_in]
18    print("E_in: ", E_in)
19    print("min E_in: ", min(E_in))
20    print("Choice: ", chr(97+max(min_E_in_index)))

```

```

1 NumbaDeprecationWarning: The 'nopython' keyword argument was not supplied to the 'numl
2 def csr_to_problem_jit(l, x_val, x_ind, x_rowptr, prob_val, prob_ind, prob_rowptr):
3 Accuracy = 100% (200/200) (classification)
4 Accuracy = 100% (200/200) (classification)
5 Accuracy = 100% (200/200) (classification)
6 Accuracy = 96% (192/200) (classification)
7 Accuracy = 76% (152/200) (classification)
8 E_in: [0.0, 0.0, 0.0, 0.04, 0.24]
9 min E_in: 0.0
10 Choice: c

```

Since $E_{in}(w_{\log_{10}(\lambda^*)=-6}) = E_{in}(w_{\log_{10}(\lambda^*)=-3}) = E_{in}(w_{\log_{10}(\lambda^*)=0}) = 0$, by selecting the largest λ , therefore, the best $\log_{10}(\lambda^*) = 0$

As a result, we should choose $[c]$ as our solution.

P14

We use the same code above, then add a `train_test_split` function and update the `main` function and the parameter of `args`. The updated functions show below.

```

1 import numpy as np
2 import random
3 from itertools import combinations_with_replacement
4 from liblinear.liblinearutil import *
5
6 def train_test_split(args, x, y):
7     rng = np.random.default_rng()
8     idx = np.arange(len(x))
9     rng.shuffle(idx)
10    x_train = x[idx[:args['split']]]
11    y_train = y[idx[:args['split']]]
12    x_test = x[idx[args['split']:]]
13    y_test = y[idx[args['split']:]]
14    return x_train, y_train, x_test, y_test

```

```

1
2 def main(args):
3     x_train, y_train = read_file(args, args['filename_train'])
4     x_train = transformation(args, x_train)
5     lamb_list = list()
6     E_val_list = list()
7     for i in range(args['repeat_time']):
8         x_train_split, y_train_split, x_test_split, y_test_split = train_test_split(a
9         args['size'] = x_train_split.shape[0]
10
11         E_val = list()
12         for lamb in args['lambda']:
13             C = 1/(2*lamb)
14             prob = problem(y_train_split, x_train_split)
15             param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))
16             m = train(prob, param)
17             p_label, p_acc, p_val = predict(y_test_split, x_test_split, m)
18             E_val.append(round(np.mean(y_test_split != p_label), 6))
19
20         min_E_val = min(E_val)
21         min_E_val_index = [i for i, v in enumerate(E_val) if v == min_E_val]
22         lamb_list.append(max(min_E_val_index))
23         E_val_list.append(min_E_val)
24
25     best_lambda = args['lambda'][max(set(lamb_list), key = lamb_list.count)]
26     print("best_lambda: ", best_lambda)

```

```

1 if __name__ == '__main__':
2     args = {
3         'dimension': 0,
4         'filename_test': "hw4_test.dat",
5         'filename_train': "hw4_train.dat",
6         'lambda': [10**(-6), 10**(-3), 10**(0), 10**3, 10**6],
7         'Q': 4,
8         'repeat_time': 256,
9         'seed': 1126,
10        'size': 0,
11        'split': 120
12    }
13
14    main(args)

```

```

1 best_lambda: 1000

```

As a result, we should choose $[d]$ as our solution.

P15

We use the same code above, then update the `main` function. The updated functions show below.

```
1  def main(args):
2      x_train, y_train = read_file(args, args['filename_train'])
3      x_train = transformation(args, x_train)
4      x_test, y_test = read_file(args, args['filename_test'])
5      x_test = transformation(args, x_test)
6      lamb_list = list()
7      E_val_list = list()
8      E_out_list = list()
9      for i in range(args['repeat_time']):
10         x_train_split, y_train_split, x_test_split, y_test_split = train_test_split(a
11         args['size'] = x_train_split.shape[0]
12
13         model_list = list()
14         E_val = list()
15         for lamb in args['lambda']:
16             C = 1/(2*lamb)
17             prob = problem(y_train_split, x_train_split)
18             param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))
19             m = train(prob, param)
20             p_label, p_acc, p_val = predict(y_test_split, x_test_split, m)
21             E_val.append(round(np.mean(y_test_split != p_label), 6))
22             model_list.append(m)
23
24         min_E_val = min(E_val)
25         min_E_val_index = [i for i, v in enumerate(E_val) if v == min_E_val]
26         lamb_list.append(max(min_E_val_index))
27         E_val_list.append(min_E_val)
28         # predict test dataset
29         p_label, p_acc, p_val = predict(y_test, x_test, model_list[max(min_E_val_inde
30         E_out_list.append(round(np.mean(y_test != p_label), 6))
31
32         best_lambda = args['lambda'][max(set(lamb_list), key = lamb_list.count)]
33         print("best_lambda: ", best_lambda)
34         print("E_out: ", np.mean(E_out_list))
```

```
1  best_lambda: 1000
2  E_out: 0.169109375
```

As a result, we should choose $[c]$ as our solution.

P16

We use the same code above, then update the `main` function. The updated functions show below.

```

1  def main(args):
2      x_train, y_train = read_file(args, args['filename_train'])
3      x_train = transformation(args, x_train)
4      x_test, y_test = read_file(args, args['filename_test'])
5      x_test = transformation(args, x_test)
6      lamb_list = list()
7      E_val_list = list()
8      E_out_list = list()
9      for i in range(args['repeat_time']):
10         x_train_split, y_train_split, x_test_split, y_test_split = train_test_split(a
11         args['size'] = x_train_split.shape[0]
12
13         model_list = list()
14         E_val = list()
15         for lamb in args['lambda']:
16             C = 1/(2*lamb)
17             prob = problem(y_train_split, x_train_split)
18             param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))
19             m = train(prob, param)
20             p_label, p_acc, p_val = predict(y_test_split, x_test_split, m)
21             E_val.append(round(np.mean(y_test_split != p_label), 6))
22             model_list.append(m)
23
24             min_E_val = min(E_val)
25             min_E_val_index = [i for i, v in enumerate(E_val) if v == min_E_val]
26             lamb_list.append(max(min_E_val_index))
27             E_val_list.append(min_E_val)
28
29             # rebuild the model from raw training dataset
30             prob = problem(y_train, x_train)
31             param = parameter('-s 0 -c {} -e 0.000001 -q'.format(1/(2*(args['lambda'])[max
32             m = train(prob, param)
33             # predict from raw testing dataset
34             p_label, p_acc, p_val = predict(y_test, x_test, m)
35             E_out_list.append(round(np.mean(y_test != p_label), 6))
36
37             best_lambda = args['lambda'][max(set(lamb_list), key = lamb_list.count)]
38             print("best_lambda: ", best_lambda)
39             print("E_out: ", np.mean(E_out_list))

```

```

1  best_lambda: 1000
2  E_out: 0.15015624999999996

```

As a result, we should choose $[b]$ as our solution.

P17

We use the same code above, then add a `KFold` function and update the `main` function and the parameter of `args`. The updated functions show below.

```

1 def KFold(args, x, y):
2     rng = np.random.default_rng()
3     idx = np.arange(len(x))
4     rng.shuffle(idx)
5     x_split = np.array([x[idx[i*args['fold_size']:(i+1)*args['fold_size']]] for i in range(args['fold'])])
6     y_split = np.array([y[idx[i*args['fold_size']:(i+1)*args['fold_size']]] for i in range(args['fold'])])
7     return x_split, y_split

```

```

1 def main(args):
2     x_train, y_train = read_file(args, args['filename_train'])
3     x_train = transformation(args, x_train)
4     x_test, y_test = read_file(args, args['filename_test'])
5     x_test = transformation(args, x_test)
6     args['size'] = x_train.shape[0]
7     args['fold_size'] = int(args['size'] / args['fold'])
8
9     E_cv_list = list()
10    for i in range(args['repeat_time']):
11        E_cv = list()
12        x_split, y_split = KFold(args, x_train, y_train)
13
14        for lamb in args['lambda']:
15            C = 1/(2*lamb)
16            E_val = list()
17            for fold in range(args['fold']):
18                x_valid_fold = x_split[fold]
19                y_valid_fold = y_split[fold]
20                x_train_fold = np.vstack(x_split[j] for j in range(args['fold']) if j != fold)
21                y_train_fold = np.hstack(y_split[j] for j in range(args['fold']) if j != fold)
22
23                prob = problem(y_train_fold, x_train_fold)
24                param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))
25                m = train(prob, param)
26                p_label, p_acc, p_val = predict(y_valid_fold, x_valid_fold, m)
27                E_val.append(round(np.mean(y_valid_fold != p_label), 6))
28            E_cv.append(np.mean(E_val))
29        E_cv_list.append(min(E_cv))
30
31    print("E_cv: ", np.mean(E_cv_list))

```



```

1  if __name__ == '__main__':
2      args = {
3          'dimension': 0,
4          'filename_test': "hw4_test.dat",
5          'filename_train': "hw4_train.dat",
6          'fold': 5,
7          'fold_size': 0, # size of each fold
8          'lambda': [10**(-6), 10**(-3), 10**(0), 10**3, 10**6],
9          'Q': 4,
10         'repeat_time': 256,
11         'seed': 1126,
12         'size': 0,
13         'split': 120
14     }
15
16     main(args)

```

```

1  E_cv: 0.12927734375

```

As a result, we should choose $[a]$ as our solution.

P18

Since L1-regularized logistic regression `-s 6` in `liblinear` solves

$$w = \min_w \sum |w_j| + C \sum_{i=1}^N \log(1 + \exp(-y_i w^T x_i))$$

and we solve

$$w_\lambda = \operatorname{argmin}_w \frac{\lambda}{N} \|w\|_1 + \frac{1}{N} \sum_{n=1}^N \ln(1 + \exp(-y_n w^T \Phi_4(x_n)))$$

Then we know that

$$1 = \frac{\lambda}{N}, C = \frac{1}{N} \Rightarrow C = \frac{1}{\lambda}.$$

The parameter `-c cost` : set the parameter `C` (default 1) should be set to $C = \frac{1}{\lambda}$.

We use the same code from `p12`, then update parameter `c` and `param` in the `main` function. The updated functions show below.

```

1  C = 1/(lamb)
2  param = parameter('-s 6 -c {} -e 0.000001 -q'.format(C))

```

```

1 NumbaDeprecationWarning: The 'nopython' keyword argument was not supplied to the 'numl
2 def csr_to_problem_jit(l, x_val, x_ind, x_rowptr, prob_val, prob_ind, prob_rowptr):
3 Accuracy = 76.8% (384/500) (classification)
4 Accuracy = 83.6% (418/500) (classification)
5 Accuracy = 84.6% (423/500) (classification)
6 Accuracy = 68% (340/500) (classification)
7 Accuracy = 49.2% (246/500) (classification)
8 E_out: [0.232, 0.164, 0.154, 0.32, 0.508]
9 min E_out: 0.154
10 Choice: c

```

As a result, we should choose `[c]` as our solution.

P19

We use the same code above, then update the `main` function. The updated functions show below.

```

1 def main(args):
2     x_train, y_train = read_file(args, args['filename_train'])
3     x_test, y_test = read_file(args, args['filename_test'])
4     x_train = transformation(args, x_train)
5     x_test = transformation(args, x_test)
6     args['dimension'] = x_train.shape[1]
7
8     model_list = list()
9     E_out = list()
10    for lamb in args['lambda']:
11        C = 1/(lamb)
12        prob = problem(y_train, x_train)
13        param = parameter('-s 6 -c {} -e 0.000001 -q'.format(C))
14        m = train(prob, param)
15        model_list.append(np.array([m.w[i] for i in range(args['dimension']) if np.ab:
16        p_label, p_acc, p_val = predict(y_test, x_test, m)
17        E_out.append(round(np.mean(y_test != p_label), 6))
18
19    min_E_out = min(E_out)
20    min_E_out_index = [i for i, v in enumerate(E_out) if v == min_E_out]
21    print("|m.w[i] <= 10**(-6)|: ", len(model_list[min_E_out_index[0]]))

```

```

1 NumbaDeprecationWarning: The 'nopython' keyword argument was not supplied to the 'numl
2 def csr_to_problem_jit(l, x_val, x_ind, x_rowptr, prob_val, prob_ind, prob_rowptr):
3 Accuracy = 76.8% (384/500) (classification)
4 Accuracy = 83.6% (418/500) (classification)
5 Accuracy = 84.6% (423/500) (classification)
6 Accuracy = 68% (340/500) (classification)
7 Accuracy = 49.2% (246/500) (classification)
8 |m.w[i] <= 10**(-6)|: 960

```

As a result, we should choose $[e]$ as our solution.

P20

We use the same code above, then update parameter `c` and `param` in the `main` function. The updated functions show below.

```

1 C = 1/(2*lamb)
2 param = parameter('-s 0 -c {} -e 0.000001 -q'.format(C))

```

```

1 NumbaDeprecationWarning: The 'nopython' keyword argument was not supplied to the 'numl
2 def csr_to_problem_jit(l, x_val, x_ind, x_rowptr, prob_val, prob_ind, prob_rowptr):
3 Accuracy = 77.4% (387/500) (classification)
4 Accuracy = 82.2% (411/500) (classification)
5 Accuracy = 84.6% (423/500) (classification)
6 Accuracy = 85.8% (429/500) (classification)
7 Accuracy = 81.2% (406/500) (classification)
8 |m.w[i] <= 10**(-6)|: 1

```

As a result, we should choose $[a]$ as our solution.